EgoQuant: documents for the judges

2026 Shenzhen FinTechathon, WeBank International Track, Topic B: Wealth Advisory Agent. Status: 20 October 2026. Live system: https://fintechaton-frontend.vercel.app · API: https://web-production-59960.up.railway.app (`/docs`) · Repositories (public copies of the frozen state, tag `submission-20oct`): https://github.com/EgoQuantAI/egoquant-backend (Python, FastAPI) and https://github.com/EgoQuantAI/egoquant-frontend (Next.js). Our private development repositories are called `fintechaton-backend` and `fintechaton-frontend`; the documents use those names in paths.

Five-minute path

1. Read the one-paragraph summary and the three promises in **01 Project summary** (sections 1 to 3).
2. Try the product with **05 User guide and demo**: the ten-minute path, a demo account, and the on-chain try-out.
3. Open **02 Security and compliance report**: the design in ten sentences (section 1) and the residual risks (section 11).
4. Skim **07 Judges' questions** for the questions we expect, with evidence.

All documents

Document	What it answers	Read it if you are
01 Project summary	What it is, the user's journey, architecture, how a portfolio is built, what is real and what is invented	everyone
02 Security and compliance report	Suitability, access control, privacy, AI governance, on-chain seal, resilience, evidence, residual risks	a security or compliance reviewer
03 Methods and models, part A	Questionnaire and scoring, the Risk Game blend, behaviour (Agent B), preferences, market data (Agent C)	a technical reviewer
03b Methods and models, part B	The category limits table, portfolio construction (Agent A), the goal simulation and its calibration, the Risk Engine, the trade check, the report, the advisor	a technical reviewer
04 Testing and evidence	Test counts by area, red-teams, backtests, probes, QA findings and how to re-run everything	an auditor
05 User guide and demo	The path through the product, with screenshots; demo accounts; the seal try-out	a judge using the system
06 Limitations and roadmap	What it is not, where it is weak, what we measured, what we would do next	everyone
07 Judges' questions	Short answers with evidence	everyone
08 Frontend and API reference	Pages, the proxy, every endpoint, security headers	a technical reviewer
09 The on-chain seal	What is sealed and where it goes, what the seal proves and does not, how to check it in two minutes	a judge interested in the blockchain part

Existing deep-dive documents (in `docs/`)

`SECURITY.md` (threat table T1-T19 and 28 gaps) · `RISK_GAME_BLEND.md` · `V2_DECISION_SHEET.md` · `ANCHORING.md` · `ONCHAIN_DECISIONS.md` · `ON_CHAIN_EVIDENCE.md` · `JUDGE_GUIDE.md` · `DEPLOY_ENV.md` · `AGENT_D_LIMITATIONS.md` ·

`egoquant/risk_profile/README.md` · `egoquant/agents/behavioral/README.md` · `egoquant/briefing/README.md` · `eval/results/regime_v2_summary.md`.

Three rules the whole system follows

1. **Numbers that matter come from deterministic, tested code; AI only explains.** (One labelled exception: Agent D's news-tone label, which never sets a number.)
2. **An independent check has the final word.** The Risk Engine re-computes every limit from the weights alone.
3. **We say what is not done.** Every claim here names a file, a test or a measured number; every weakness is listed.

Everything is hypothetical and based on past data. Nothing here is a forecast or a recommendation to invest. EgoQuant places no orders and moves no money.